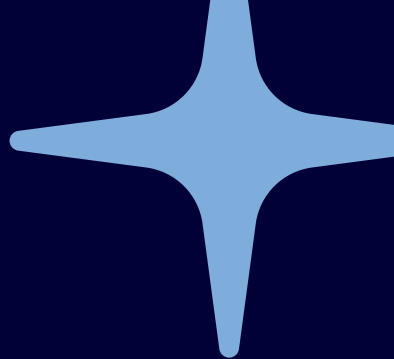


Low Rank Adapters (LoRA)

Advanced LLM Night

Johannes Kohlmann – Findable

January 15th, 2025



WHO IS THIS GUY?



WHO IS THIS GUY?



Johannes Kohlmann (Me)



WHO IS THIS GUY?

- Originally from a tiny village in southern Germany



Johannes Kohlmann (Me)



WHO IS THIS GUY?

- Originally from a tiny village in southern Germany
- Masters degree in Computer Science (ML & theoretical CS focus)



Johannes Kohlmann (Me)



WHO IS THIS GUY?

- Originally from a tiny village in southern Germany
- Masters degree in Computer Science (ML & theoretical CS focus)
- Moved to Oslo in the beginning of 2024



Johannes Kohlmann (Me)



WHO IS THIS GUY?

- Originally from a tiny village in southern Germany
- Masters degree in Computer Science (ML & theoretical CS focus)
- Moved to Oslo in the beginning of 2024
- Software Engineer at Findable since June 2024



Johannes Kohlmann (Me)



HOW DO MODERN LLMs WORK?



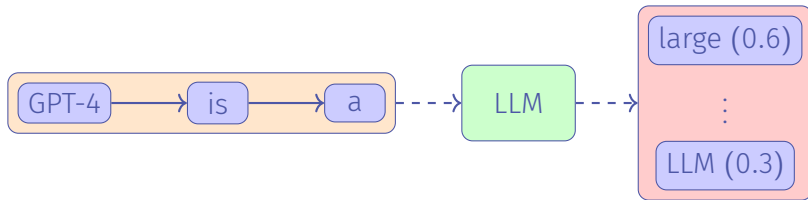
HOW DO MODERN LLMs WORK?

Modern LLMs predict the most probable word given the input:




HOW DO MODERN LLMs WORK?

Modern LLMs predict the most probable word given the input:



HOW DO MODERN LLMs WORK?

Modern LLMs predict the most probable word given the input:

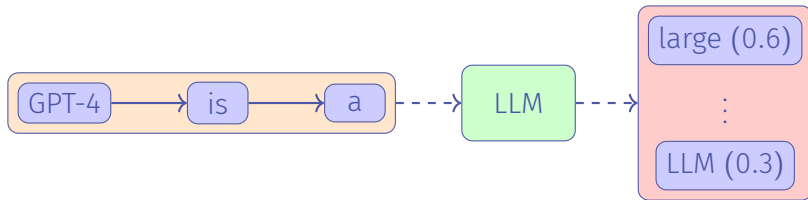


Sooo.... how do they learn to do that?



HOW DO MODERN LLMs WORK?

Modern LLMs predict the most probable word given the input:



Sooo.... how do they learn to do that?

Next Token Prediction

Given a dataset \mathcal{D} of sentences s and target words s_t , minimize the loss \mathcal{L}

$$\mathcal{L}(\mathcal{D}) = \sum_{s \in \mathcal{D}} L(\text{LLM}(s), s_t)$$



PRE-TRAINING DONE. NOW WHAT?



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).

However, to utilize them for specific tasks, we need to **finetune**!



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).

However, to utilize them for specific tasks, we need to **finetune**!

Finetuning

Finetuning specializes a pre-trained model for a specific task. This is often done by changing/adding weights.



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).

However, to utilize them for specific tasks, we need to **finetune**!

Finetuning

Finetuning specializes a pre-trained model for a specific task. This is often done by changing/adding weights.

There is a variety of approaches to finetuning:



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).

However, to utilize them for specific tasks, we need to **finetune**!

Finetuning

Finetuning specializes a pre-trained model for a specific task. This is often done by changing/adding weights.

There is a variety of approaches to finetuning:

- Full Model Finetuning (Expensive, Hard to maintain)



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).

However, to utilize them for specific tasks, we need to **finetune**!

Finetuning

Finetuning specializes a pre-trained model for a specific task. This is often done by changing/adding weights.

There is a variety of approaches to finetuning:

- Full Model Finetuning (Expensive, Hard to maintain)
- Additional Adapter Layers (Limited capacity, Inference latency)



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).

However, to utilize them for specific tasks, we need to **finetune**!

Finetuning

Finetuning specializes a pre-trained model for a specific task. This is often done by changing/adding weights.

There is a variety of approaches to finetuning:

- Full Model Finetuning (Expensive, Hard to maintain)
- Additional Adapter Layers (Limited capacity, Inference latency)
- Prompt Optimization (Limited capacity & generalization)



PRE-TRAINING DONE. NOW WHAT?

After pre-training, LLMs are generalists (a.k.a. Foundational Models).

However, to utilize them for specific tasks, we need to **finetune**!

Finetuning

Finetuning specializes a pre-trained model for a specific task. This is often done by changing/adding weights.

There is a variety of approaches to finetuning:

- Full Model Finetuning (Expensive, Hard to maintain)
- Additional Adapter Layers (Limited capacity, Inference latency)
- Prompt Optimization (Limited capacity & generalization)

Sooo... what now?



FULL MODEL FINETUNING



FULL MODEL FINETUNING

During full model finetuning, we do small updates to all weights of the model:



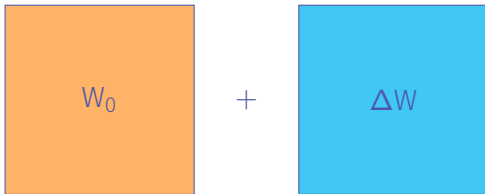
FULL MODEL FINETUNING

During full model finetuning, we do small updates to all weights of the model:



FULL MODEL FINETUNING

During full model finetuning, we do small updates to all weights of the model:



The diagram illustrates the process of full model finetuning. It consists of two colored squares, one orange and one blue, separated by a plus sign. The orange square contains the text W_0 , representing the initial weight matrix. The blue square contains the text ΔW , representing the small update matrix. The plus sign indicates that the update matrix is added to the initial weight matrix.

$$W_0 + \Delta W$$

FULL MODEL FINETUNING

During full model finetuning, we do small updates to all weights of the model:



FULL MODEL FINETUNING

During full model finetuning, we do small updates to all weights of the model:



Therefore, finetuning a single fully-connected layer has the same complexity as learning it from scratch!



FULL MODEL FINETUNING

During full model finetuning, we do small updates to all weights of the model:



Therefore, finetuning a single fully-connected layer has the same complexity as learning it from scratch!

So, doing this for all layers makes finetuning as compute-intensive as the original pre-training!



A LITTLE BIT OF MATH



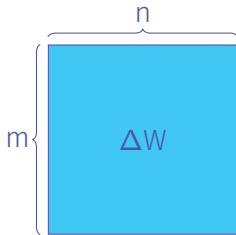
A LITTLE BIT OF MATH

How can we reduce the number of trainable parameters during finetuning?



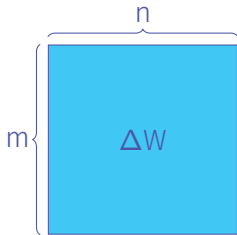
A LITTLE BIT OF MATH

How can we reduce the number of trainable parameters during finetuning?



A LITTLE BIT OF MATH

How can we reduce the number of trainable parameters during finetuning?



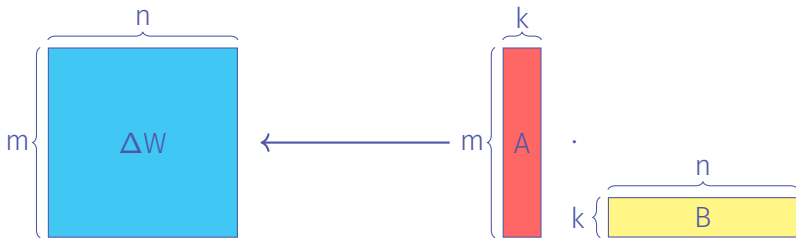
Matrix multiplication "eliminates" the "inner" dimension:

$$(m \times n) \longleftarrow (m \times k) \cdot (k \times n)$$



A LITTLE BIT OF MATH

How can we reduce the number of trainable parameters during finetuning?




Matrix multiplication "eliminates" the "inner" dimension:

$$(m \times n) \longleftarrow (m \times k) \cdot (k \times n)$$



A LITTLE BIT OF MATH

How can we reduce the number of trainable parameters during finetuning?



Matrix multiplication "eliminates" the "inner" dimension:

$$(m \times n) \leftarrow (m \times k) \cdot (k \times n)$$

Therefore, we can represent ΔW by two much smaller matrices!



LoRA FOR EFFICIENT FINETUNING



LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:



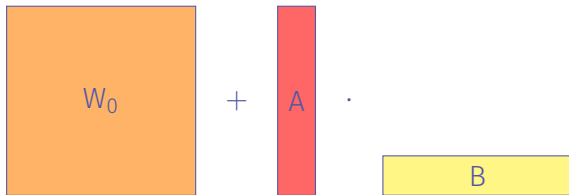
LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:




LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:


$$W_0 + A \cdot B$$

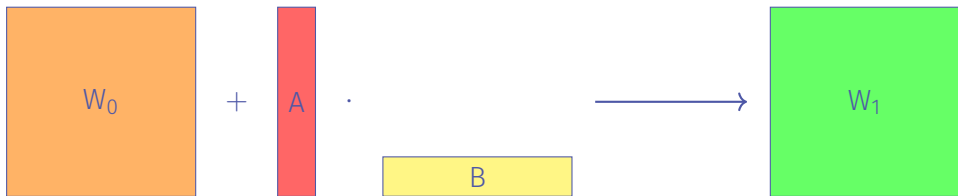
LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:



LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:

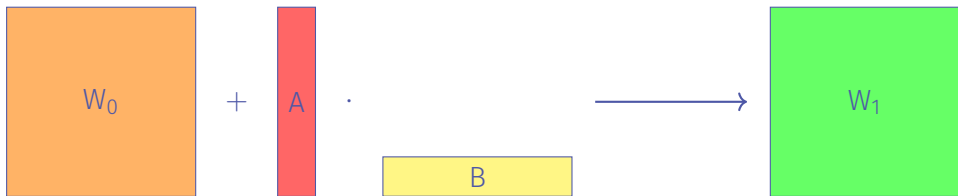


Or a little more mathy:



LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:




Or a little more mathy:

W_0



LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:



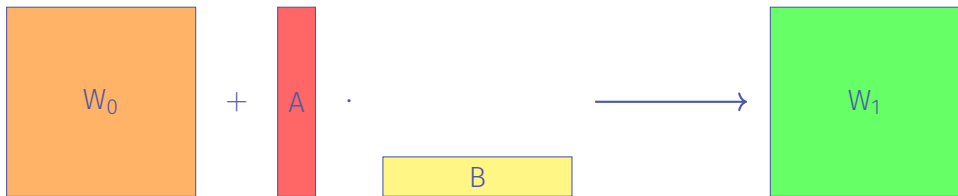
Or a little more mathy:

$$W_0 + A \cdot B$$



LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:




Or a little more mathy:

$$W_0 + A \cdot B \approx W_0 + \Delta W$$



LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:



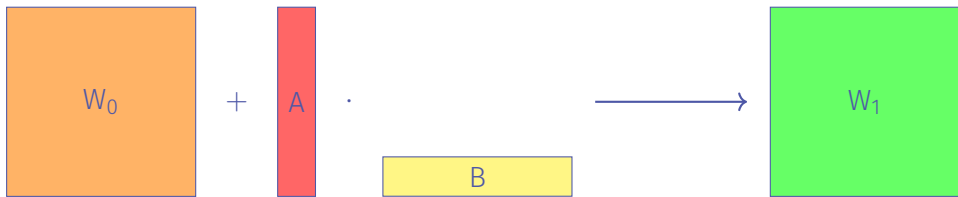
Or a little more mathy:

$$W_0 + A \cdot B \approx W_0 + \Delta W \longrightarrow W_1$$



LoRA FOR EFFICIENT FINETUNING

Low-rank Adapters (LoRA) make use of this trick:



Or a little more mathy:

$$W_0 + A \cdot B \approx W_0 + \Delta W \longrightarrow W_1$$

Parameter reduction from $m \cdot n$ to $k \cdot (m + n)$ which is very significant for $m, n \gg k$!



CLASSIFYING MNIST



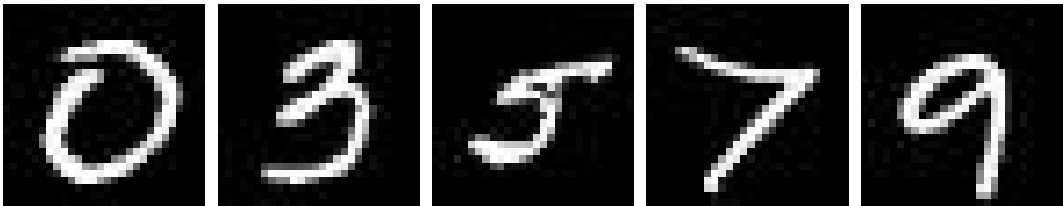
CLASSIFYING MNIST

Imagine, we are tasked to recognize hand-written digits:



CLASSIFYING MNIST

Imagine, we are tasked to recognize hand-written digits:




Example images from the MNIST dataset



CLASSIFYING MNIST

Imagine, we are tasked to recognize hand-written digits:



Example images from the MNIST dataset

Task: Given an input image, determine the digit this image depicts.



MODEL ARCHITECTURE



MODEL ARCHITECTURE

We'll use a simple, fully-connected model:



MODEL ARCHITECTURE

We'll use a simple, fully-connected model:

```
class SimpleMnistClassifier(nn.Module):  
    def __init__(self):  
        super(SimpleMnistClassifier, self).__init__()  
        self._fc1 = nn.Linear(784, 256) # 28 * 28 = 784  
        self._fc2 = nn.Linear(256, 128)  
        self._fc3 = nn.Linear(128, 10)  
  
    def forward(x: torch.Tensor) -> torch.Tensor:  
        # Uses ReLU and returns logits
```



MODEL ARCHITECTURE

We'll use a simple, fully-connected model:

```
class SimpleMnistClassifier(nn.Module):  
    def __init__(self):  
        super(SimpleMnistClassifier, self).__init__()  
        self._fc1 = nn.Linear(784, 256) # 28 * 28 = 784  
        self._fc2 = nn.Linear(256, 128)  
        self._fc3 = nn.Linear(128, 10)  
  
    def forward(x: torch.Tensor) -> torch.Tensor:  
        # Uses ReLU and returns logits
```

This model has 235.146 trainable parameters and reaches a test accuracy of 97.78%.



FLIPPED MNIST




FLIPPED MNIST

It turns out, that we did not train on exactly the right data:



FLIPPED MNIST

It turns out, that we did not train on exactly the right data:

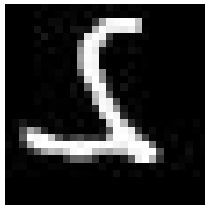


Flipped-MNIST (F-MNIST) has 30% of its two's and three's flipped



FLIPPED MNIST

It turns out, that we did not train on exactly the right data:



Flipped-MNIST (F-MNIST) has 30% of its two's and three's flipped

How does our model perform on this data? Is it going to be able to generalize by itself?



EVALUATION ON F-MNIST





EVALUATION ON F-MNIST

Evaluating our model on F-MNIST yields a test accuracy of 92.87%.



EVALUATION ON F-MNIST

Evaluating our model on F-MNIST yields a test accuracy of 92.87%.



Confusion matrices for MNIST (left) and F-MNIST (right)



SALVAGING OUR EFFORT WITH LoRA



SALVAGING OUR EFFORT WITH LoRA

We can define a LoRA layer as follows:



SALVAGING OUR EFFORT WITH LoRA

We can define a LoRA layer as follows:

```
class LoRALayer(nn.Module):  
    def __init__(self, inner: nn.Linear, k: int):  
        self._inner = inner  
        self._inner.requires_grad_(False)
```



SALVAGING OUR EFFORT WITH LORA

We can define a LoRA layer as follows:

```
class LoRALayer(nn.Module):  
    def __init__(self, inner: nn.Linear, k: int):  
        self._inner = inner  
        self._inner.requires_grad_(False)  
        n, m = inner.in_features, inner.out_features  
        self._A = nn.Parameter(torch.randn(n, k))  
        self._B = nn.Parameter(torch.zeros(k, m))
```



SALVAGING OUR EFFORT WITH LORA

We can define a LoRA layer as follows:

```
class LoRALayer(nn.Module):  
    def __init__(self, inner: nn.Linear, k: int):  
        self._inner = inner  
        self._inner.requires_grad_(False)  
        n, m = inner.in_features, inner.out_features  
        self._A = nn.Parameter(torch.randn(n, k))  
        self._B = nn.Parameter(torch.zeros(k, m))  
  
    def forward(self, x: torch.Tensor) -> torch.Tensor:  
        return self._inner(x) + x @ self._A @ self._B
```



RESULTS WITH LoRA



RESULTS WITH LoRA

Augmenting the linear layers with LoRA ($k = 4$) yields a model with only

$$(784 \cdot 4 + 4 \cdot 256) + (256 \cdot 4 + 4 \cdot 128) + (128 \cdot 4 + 4 \cdot 10) = 6248$$

trainable parameters.



RESULTS WITH LoRA

Augmenting the linear layers with LoRA ($k = 4$) yields a model with only

$$(784 \cdot 4 + 4 \cdot 256) + (256 \cdot 4 + 4 \cdot 128) + (128 \cdot 4 + 4 \cdot 10) = 6248$$

trainable parameters. That is only 3% of the original parameters!



RESULTS WITH LoRA

Augmenting the linear layers with LoRA ($k = 4$) yields a model with only

$$(784 \cdot 4 + 4 \cdot 256) + (256 \cdot 4 + 4 \cdot 128) + (128 \cdot 4 + 4 \cdot 10) = 6248$$

trainable parameters. That is only **3% of the original parameters!**

The finetuned model reaches a test accuracy of 97.74%, i.e. it is as good as the original model was on MNIST!



RESULTS WITH LoRA

Augmenting the linear layers with LoRA ($k = 4$) yields a model with only

$$(784 \cdot 4 + 4 \cdot 256) + (256 \cdot 4 + 4 \cdot 128) + (128 \cdot 4 + 4 \cdot 10) = 6248$$

trainable parameters. That is only **3% of the original parameters!**

The finetuned model reaches a test accuracy of 97.74%, i.e. it is as good as the original model was on MNIST!

Imagine doing this with GPT3-175B! Checkpoint sizes are **reduced by 99.99%**, finetuning is **sped up by 25%** and you only need **a third of the GPUs to train!**



CONCLUSION

- Based on the LoRA paper by Hu et al.
- Code and slides will be available online.
- Thank you for listening! Questions?

